


SigRank — Build Archetype Analytics


When the field is filtered by composition type, the yield spreads 9,000x and leverage spreads 133x · 1,586 operators (HCM) · Generated 2026-08-10




Archetype Distribution



Leverage & Velocity




Token Composition (4-Pillar Breakdown)




Activity Level



Signal Quality



Top Platform per Archetype



Full Archetype Data Table

#	NAME	FAMILY	N	%	YIELD Y	LEVERAGE	VELOCITY	SNR	TOKENS (MED)	TOK/DAY	PLATFORM	COMPOSITION
0	CONVERGENT	Convergence	104	6.6%	257.30	245.3x	1.049	0.5119	2.30B	46.5M	anthropic	
1	KINETIC	Generation	113	7.1%	469.35	361.5x	1.520	0.6032	3.82B	72.1M	anthropic	
2	INPUT-BOUND	Reuse Depth	108	6.8%	0.05	2.7x	0.019	0.0191	2.53B	36.9M	openai	
3	PRIMING	Reuse Depth	149	9.4%	0.32	7.7x	0.042	0.0407	3.07B	40.6M	openai	
4	CONTEXTUAL	Reuse Depth	186	11.7%	0.84	12.7x	0.066	0.0618	4.55B	49.9M	openai	
5	DEEP READER	Reuse Depth	169	10.7%	1.52	18.4x	0.081	0.0750	7.96B	95.7M	openai	
6	ARCHIVIST	Reuse Depth	186	11.7%	3.54	29.8x	0.124	0.1100	16.63B	172.8M	openai	
7	BUILDER	Active Construction	278	17.5%	1.12	13.7x	0.077	0.0715	4.01B	45.0M	anthropic	
8	RECURSIVE	Active Construction	132	8.3%	7.23	38.9x	0.191	0.1602	6.69B	79.6M	anthropic	
9	AMPLIFIER	Active Construction	161	10.2%	26.53	77.2x	0.329	0.2474	6.55B	92.7M	anthropic	

Methodology

Method: deterministic_build_archetypes — no clustering, no randomness

Classification: Each operator is classified into one of 10 build archetypes based on three derived dimensions of their token cascade:
Leverage = cache_read / input (how much you reuse vs fresh input)
Velocity = output / input (how much you generate vs take in)
Construction = cache_write / cache_read (how much new context you build per read)

Precedence (first match wins):
1. **CONVERGENT** — P80 on all 3 axes (leverage ≥ 74.6, velocity ≥ 0.34, construction ≥ 0.043)
2. **KINETIC** — velocity ≥ 0.80
3. **Construction branch** — construction ≥ 0.02 (BUILDER / RECURSIVE / AMPLIFIER by leverage)
4. **Reuse depth branch** — else (INPUT-BOUND / PRIMING / CONTEXTUAL / DEEP READER / ARCHIVIST by leverage)

Source: public/data/archetypes.json generated by scripts/gen-build-archetypes.mjs from public/data/field-analysis.json
Population: human_center_of_mass (1,627 total, 1,586 HCM, outliers excluded)
Replaces: Previous 8 K-Means archetypes (silhouette 0.544/0.377) — now deterministic with clear definitions

Citation
McHenry, D. J. (2026). SigRank Two-Axis Operator Taxonomy: Finalized Datasets and Analytics Dashboards (v3.0) [Dataset]. Zenodo. <https://doi.org/10.5281/zenodo.21876660>

Concept DOI (resolves to latest version)
<https://doi.org/10.5281/zenodo.21876675>

Author
Deric J. McHenry · ORCID 0009-0002-9904-5390
SigRank · signalaf.com · [@burnmydays](https://twitter.com/burnmydays)

License
Creative Commons Attribution 4.0 International (CC-BY 4.0)

BibTeX

```
@dataset{mchenry2026sigrank,
  author      = {McHenry, Deric J.},
  title       = {{SigRank Two-Axis Operator Taxonomy: Finalized Datasets and Analytics Dashboards (v3.0)}},
  year        = {2026},
  publisher    = {Zenodo},
  version     = {3.0},
  doi         = {10.5281/zenodo.21876660},
  url         = {https://doi.org/10.5281/zenodo.21876660}}
```

SigRank — Experience Ladder Analytics

When the field is filtered by total tokens, the experience span ranges from 0 to 84 billion · 1,626 operators · 24 equal-population stages · Generated 2026-08-10


EXPERIENCE SPAN
84B tokens
IGNITER III (0) → ARCH+ I floor (84.19B)

STAGE WIDTH
~68 ops
Equal-population calibration


TIERS
8 + 1
8 base + TRANSMITTER peak badge

STAGES
24
8 tiers × 3 sub-stages (I/II/III)


Tier Distribution




Token Ranges



Stage Distribution by Tier



Tier Population



Full 24-Stage Data Table

STAGE	TIER	FLOOR	CEILING	OPERATORS	% OF FIELD
ARCH+ I	ARCH+	84.19B	7.07T	67	4.1%
ARCH+ II	ARCH+	46.53B	83.80B	68	4.2%
ARCH+ III	ARCH+	33.12B	46.01B	68	4.2%
ARCH I	ARCH	23.90B	33.11B	68	4.2%
ARCH II	ARCH	18.40B	23.89B	67	4.1%
ARCH III	ARCH	15.53B	18.39B	68	4.2%
POWER I	POWER	12.27B	15.51B	68	4.2%
POWER II	POWER	10.22B	12.11B	68	4.2%
POWER III	POWER	8.77B	10.19B	67	4.1%
BASE I	BASE	7.08B	8.67B	68	4.2%
BASE II	BASE	5.89B	7.07B	68	4.2%
BASE III	BASE	4.91B	5.88B	68	4.2%
SEEKER I	SEEKER	4.15B	4.91B	67	4.1%
SEEKER II	SEEKER	3.59B	4.15B	68	4.2%
SEEKER III	SEEKER	2.98B	3.59B	68	4.2%
REFINER I	REFINER	2.50B	2.96B	68	4.2%
REFINER II	REFINER	2.01B	2.50B	67	4.1%
REFINER III	REFINER	1.55B	1.99B	68	4.2%
BEARER I	BEARER	1.18B	1.54B	68	4.2%
BEARER II	BEARER	859M	1.17B	68	4.2%
BEARER III	BEARER	640M	859M	67	4.1%
IGNITER I	IGNITER	305M	639M	68	4.2%
IGNITER II	IGNITER	132M	305M	68	4.2%
IGNITER III	IGNITER	123K	131M	68	4.2%

Methodology

Method: equal_population_24_stages — Option C calibration

Classification: Each operator is assigned to one of 24 experience stages based on their **total tokens** (input + output + cacheCreate + cacheRead). The ladder tracks how much real work an operator has done — not how efficient that work is.

Tier structure: 8 base tiers, each split into 3 sub-stages (I/II/III):
ARCH+ (K.01) — deepest field experience · volume that became architecture
ARCH (K.02) — system builders · sustained volume, coherent output
POWER (K.03) — above the center · volume compounding
BASE (K.04) — the center of the field · the average operator's experience
SEEKER (K.05) — approaching the center · experience accumulating
REFINER (K.06) — practicing with purpose · early sustained volume
BEARER (K.07) — quiet accumulation · the first real volume
IGNITER (K.08) — dormant potential · the still soul, waiting

TRANSMITTER is not on the ladder. It is a temporary peak badge (RS.08) that any tier can earn during a high-frequency, high-resonance window. It lapses when frequency or resonance drops.

Source: public/data/class-distribution-reference.json from public/data/field-analysis.json (1,626 operators, HCM cut)
Calibration: Equal-population (Option C) — each stage holds ~68 operators. Floor = minimum total tokens; ceiling = next stage's floor.
Replaces: Previous yield-based 4-tier system (IGNITER/SEEKER/BUILDER/TRANSMITTER by yield range)

Citation
McHenry, D. J. (2026). SigRank Two-Axis Operator Taxonomy: Finalized Datasets and Analytics Dashboards (v3.0) [Dataset]. Zenodo. <https://doi.org/10.5281/zenodo.21876660>

Concept DOI (resolves to latest version)
<https://doi.org/10.5281/zenodo.21875675>

Author
Deric J. McHenry · ORCID 0009-0002-9904-5390
SigRank · signalaf.com · MO\$ES™ · @burnmydays


License
Creative Commons Attribution 4.0 International (CC-BY 4.0)

BibTeX

```
@dataset{mchenry2026sigrank,
  author      = {McHenry, Deric J.},
  title       = {(SigRank Two-Axis Operator Taxonomy: Finalized Datasets and Analytics Dashboards (v3.0))},
  year        = {2026},
  publisher    = {Zenodo},
  version     = {3.0},
  doi         = {10.5281/zenodo.21876660},
  url         = {https://doi.org/10.5281/zenodo.21876660}}
```


Cross-Metric Analytics

Divergences & convergences across the 10 build archetypes · multi-metric combined charts · 1,586 HCM operators · 2026-08-10




KEY DIVERGENCE
The field splits on **leverage**, not velocity. 6 of 10 archetypes have velocity between 0.04 and 0.12 — they all generate at similar rates. What separates them is how much they reuse: **INPUT-BOUND** reads 2.7x its input in cache; **ARCHIVIST** reads 30x; **KINETIC** reads 362x. Yield compounds this gap multiplicatively, producing a 4,694x spread from the same basic operating tempo.


1. Leverage vs Yield — the compounding scatter




2. Archetype radar — four workflow signatures



3. Yield bars + Velocity line — where velocity diverges from yield




4. Construction vs Leverage — the build-to-reuse pipeline



KEY CONVERGENCE
Cache-read percentage converges across the entire field: 88–96% of all tokens are cache reads for every archetype. The field is **unanimously cache-read dominant**. The divergence is in what feeds those reads — construction ratio. The Active Construction family writes 3–6% of their token budget as new cache. The Reuse Depth family writes 0.2–0.3%. Same read behavior, completely different write behavior. That's where the archetype signal lives.


5. Token composition (stacked) + SNR overlay



6. Metric heatmap — archetype z-scores across 6 dimensions




7. Tokens/day vs Yield — the throughput paradox



8. Family radar — the four workflow families



9. Population vs Yield — where the field concentrates



THE POWER LAW
The top 2 archetypes (KINETIC + CONVERGENT = 217 operators, 13.7% of the field) have a combined median yield of 363. The bottom 4 (INPUT-BOUND + PRIMING + CONTEXTUAL + BUILDER = 721 operators, 45.5% of the field) have a combined median yield of 0.6. That's a 605x gap between the top eighth and the bottom half. The field is not normally distributed — it's a power law, and the archetype system makes the structure visible.

Methodology

Charts: All charts use per-archetype metrics from `archetypes.json` (10 deterministic build archetypes, 1,586 HCM operators). Radar charts normalize each metric to the field maximum (0–1 scale). The heatmap uses z-scores (standard deviations from the field mean). Scatter plots use log-log axes where indicated.

Metrics: Yield ($Y = \text{cache_read} \times \text{output} / \text{input}^2$), Leverage ($\text{cache_read} / \text{input}$), Velocity ($\text{output} / \text{input}$), SNR ($\text{output} / \text{total}$), Construction ($\text{cache_write} / \text{cache_read}$), Input share ($\text{input} / \text{total}$ tokens).

Population: human_center_of_mass (1,627 total, 1,586 HCM, outliers excluded). Input/total ratio analysis separates extreme humans from replay outliers.

Source: `public/data/archetypes.json` generated by `scripts/gen-build-archetypes.mjs` from `public/data/field-analysis.json`

Citation
McHenry, D. J. (2026). SigRank Two-Axis Operator Taxonomy: Finalized Datasets and Analytics Dashboards (v3.0) [Dataset]. Zenodo. <https://doi.org/10.5281/zenodo.21876660>

Concept DOI (resolves to latest version)
<https://doi.org/10.5281/zenodo.21875675>

Author
Deric J. McHenry · ORCID 0009-0002-9904-5390
SigRank · signalaf.com · MOSES™ · @burnmydays


License
Creative Commons Attribution 4.0 International (CC-BY 4.0)

BibTeX

```
@dataset{mcHenry2026sigrank,
  author      = {McHenry, Deric J.},
  title       = {{{SigRank Two-Axis Operator Taxonomy: Finalized Datasets and Analytics Dashboards (v3.0)}}},
  year        = {2026},
  publisher    = {Zenodo},
  version     = {3.0},
  doi         = {10.5281/zenodo.21876660},
  url        = {https://doi.org/10.5281/zenodo.21876660}
}
```


Operator Scatter

1,598 individual operators splattered across metric pairs - each dot is one human - colored by platform - 2026-08-10




READING THESE CHARTS
Every dot is a single operator. Axes are log-scaled (except SNR). The field forms **shapes** — clouds, streaks, clusters — that archetype averages hide. Look for the diagonal lines (power-law relationships), the empty quadrants (impossible combinations), and the outliers that break the pattern.


1. Yield vs Leverage — the compounding cloud



2. Yield vs Total Tokens — the volume lie




3. Velocity vs Leverage — the operating plane




THE EMPTY QUADRANT
The upper-right of the Velocity vs Leverage chart is **nearly empty**. Almost no operator has both high velocity (>10) AND high leverage (>100x). This is the tradeoff that defines the field: you either generate fast (burning input) or reuse deeply (reading cache). CONVERGENT operators are the ones who break this tradeoff — and there are only ~100 of them in 1,598.


4. SNR vs Leverage — does reuse improve signal?




5. Input Share vs Yield — the input tax




6. Cache Read % vs Cache Write % — the reuse architecture




7. Tokens/day vs Yield — throughput doesn't predict yield



8. Active Days vs Yield — does time predict skill?



9. Velocity vs SNR — signal speed vs signal purity



Methodology

Data: 1,598 operators from `public/data/field-analysis.json` with positive yield AND positive leverage (required for log-scale axes). 29 operators with yield=0 excluded from log-scale charts.

Anonymization: No handles or display names shown. Each dot is one operator identified only by platform color.

Metrics: Yield ($Y = \text{cache_read} \times \text{output} / \text{input}^2$), Leverage ($\text{cache_read} / \text{input}$), Velocity ($\text{output} / \text{input}$), SNR ($\text{output} / \text{total}$), Input Share ($\text{input} / \text{total} \times 100$), Cache Read % ($\text{cache_read} / \text{total} \times 100$), Cache Write % ($\text{cache_write} / \text{total} \times 100$).

Axes: All scatter plots use log-log scale except SNR (linear 0–1) and Active Days (linear 0–365). Point size is uniform (2px radius) to show density — overlapping dots create natural density visualization.

Source: `public/data/field-analysis.json` from `public/data/board-all_time.json` - scraped from public AI coding agent leaderboards - 2026-07-13

Citation
McHenry, D. J. (2026). SigRank Two-Axis Operator Taxonomy: Finalized Datasets and Analytics Dashboards (v3.0) [Dataset]. Zenodo. <https://doi.org/10.5281/zenodo.21876660>

Concept DOI (resolves to latest version)
<https://doi.org/10.5281/zenodo.21876675>

Author
Deric J. McHenry · ORCID 0009-0002-9904-5390
SigRank · signalaf.com · MOSES™ · @burmmydays

License
Creative Commons Attribution 4.0 International (CC-BY 4.0)

BibTeX

```
@dataset{mchenry2026sigrank,
author    = {McHenry, Deric J.},
title     = {(SigRank Two-Axis Operator Taxonomy: Finalized Datasets and Analytics Dashboards (v3.0))},
year      = {2026},
publisher = {Zenodo},
version   = {(3.0)},
doi       = {(10.5281/zenodo.21876660)},
url       = {(https://doi.org/10.5281/zenodo.21876660)}}
```

SigRank Platform Board — 17 Platforms

Source: raw-data.csv platform_breakdown · 1628 operators · CC-BY 4.0 · 2026-08-11

Sorted by total platform tokens (descending). Each platform counts every operator who uses it, not just primary.

Platform Raw — Token Volume & Cost

Platform	Ops	Volume		Tokens Sum	Median	P25	P75	Cost Sum	Cost Median	Models Med
other	869	9007.65T		9007645991382800	12.32M	557.3K	116.19M	\$1.32M	\$1.02	2
anthropic	1477	28.34T		28338885644212	1.84B	313.31M	7.02B	\$265.62M	\$1.56K	6
openai	1427	17.94T		17943907752167	1.50B	231.19M	7.09B	\$24.38M	\$886.06	6
unknown	584	6.58T		6575986322775	6.63M	246.5K	60.34M	\$16.84M	\$0.00	1
zhipu	701	483.68B		483684504483	45.53M	3.77M	281.59M	\$304.31K	\$11.09	2
deepseek	421	353.94B		353943132599	60.15M	4.06M	344.61M	\$2.18M	\$1.60	2
google	1010	304.31B		304310947658	29.47M	3.32M	157.71M	\$214.33K	\$13.11	4
minimax	649	264.05B		264052041898	18.61M	1.95M	133.83M	\$26.02K	\$0.80	2
moonshot	573	247.81B		247810814878	38.54M	5.39M	232.95M	\$4.33M	\$10.19	2
xiaomi	256	110.64B		110636587115	29.36M	3.65M	196.44M	\$2.19M	\$0.73	2
alibaba	394	90.35B		90347812379	4.58M	396.7K	40.15M	\$39.77K	\$0.56	2
xai	405	70.91B		70912966232	5.44M	653.0K	42.97M	\$7.29K	\$0.77	1
nvidia	169	2.89B		2891822242	1.39M	224.2K	6.34M	\$623.52	\$0.02	1
mistral	44	1.28B		1284108882	428.7K	51.6K	5.59M	\$381.12	\$0.09	1
bytedance	11	92.71M		92708642	913.3K	653.8K	11.00M	\$3.12	\$0.00	1
cohere	1	5.03M		5026492	5.03M	5.03M	5.03M	\$1.17	\$1.17	2
tencent	1	2.82M		2819606	2.82M	2.82M	2.82M	\$0.00	\$0.00	1

Platform Raw — Operator Aggregate Stats (Median)

Median values for operators who use each platform.

Platform	Ops	Total Med	Input Med	Output Med	Cache Read Med	Cache Write Med	Reasoning Med	Days Med	Sessions Med
other	869	6.36B	400.98M	30.17M	5.66B	74.24M	3.56M	104	371
anthropic	1477	5.69B	228.88M	25.32M	5.04B	87.63M	2.16M	90	259
openai	1427	5.85B	260.17M	26.20M	5.33B	80.12M	2.69M	92	253
unknown	584	11.20B	610.46M	53.17M	9.93B	155.79M	6.19M	133	1897
zhipu	701	7.17B	490.13M	33.84M	6.00B	84.27M	3.95M	107	519
deepseek	421	6.08B	476.29M	28.52M	4.92B	37.68M	4.12M	102	880
google	1010	7.07B	349.72M	32.17M	6.23B	97.97M	3.56M	103	280.5
minimax	649	6.32B	439.66M	29.66M	5.52B	74.24M	4.09M	106	464
moonshot	573	6.55B	507.53M	30.48M	5.63B	75.44M	4.37M	112	570
xiaomi	256	5.47B	473.74M	27.32M	4.62B	32.02M	4.00M	103	936.5
alibaba	394	6.80B	562.50M	32.68M	5.94B	68.43M	4.19M	108	633
xai	405	9.89B	602.54M	45.38M	8.57B	105.64M	5.98M	127	684
nvidia	169	7.76B	536.41M	34.23M	6.15B	35.87M	5.24M	116	1282
mistral	44	6.50B	667.15M	24.49M	5.18B	32.70M	3.21M	106	300.5
bytedance	11	7.81B	507.57M	33.60M	5.33B	97.30M	2.41M	84	725
cohere	1	13.54B	2.62B	48.45M	10.53B	337.22M	4.74M	122	2159
tencent	1	37.63B	14.44B	200.32M	22.93B	13.00M	49.94M	161	268

Platform Metrics — Derived Analytical

Median yield, SNR, leverage, velocity, compression per platform.

Platform	Ops	Yield Med	Yield P25	Yield P75	SNR Med	Leverage Med	Velocity Med	Compress Med	Tok/Day Med
other	869	1.151	0.298	3.600	0.0692	14.40	0.07	0.985	64.74M
anthropic	1477	2.177	0.657	15.906	0.0947	22.75	0.10	0.979	69.63M
openai	1427	1.678	0.546	7.247	0.0827	18.51	0.09	0.984	69.00M
unknown	584	1.617	0.461	4.946	0.0833	17.41	0.09	0.984	90.36M
zhipu	701	0.969	0.239	3.180	0.0661	13.59	0.07	0.986	75.36M
deepseek	421	0.827	0.194	2.080	0.0602	12.53	0.06	0.992	61.45M
google	1010	1.379	0.374	5.251	0.0752	16.64	0.08	0.984	76.94M
minimax	649	0.883	0.220	2.625	0.0632	12.65	0.07	0.986	64.85M
moonshot	573	0.843	0.211	2.450	0.0619	12.30	0.07	0.986	65.52M
xiaomi	256	0.599	0.163	1.847	0.0570	10.22	0.06	0.991	54.84M
alibaba	394	0.814	0.191	2.711	0.0613	12.20	0.07	0.987	76.28M
xai	405	1.062	0.259	3.109	0.0687	14.55	0.07	0.985	81.22M
nvidia	169	0.547	0.145	1.725	0.0510	9.60	0.05	0.993	65.52M
mistral	44	0.411	0.131	1.436	0.0402	8.74	0.04	0.993	81.39M
bytedance	11	0.308	0.060	1.089	0.0476	7.94	0.05	0.991	62.98M
cohere	1	0.074	0.074	0.074	0.0182	4.02	0.02	0.969	111.02M
tencent	1	0.022	0.022	0.022	0.0137	1.59	0.01	0.999	233.71M

DOI: 10.5281/zenodo.21876660

Zenodo

v3.0

CC-BY 4.0

ORCID 0009-0002-9904-5390

SigRank

MOSES™

McHenry, D. J. (2026). SigRank Two-Axis Operator Taxonomy: Finalized Datasets and Analytics Dashboards (v3.0) [Dataset]. Zenodo. <https://doi.org/10.5281/zenodo.21876660>

SigRank Model Board — 3304 Models (top 100 shown)

Source: raw-data.csv model_usage · 1628 operators · CC-BY 4.0 · 2026-08-11

Model Raw — By Total Token Volume (top 100)

MODEL		OPS	VOLUME	TOKENS SUM	MEDIAN	P25	P75	COST SUM	COST MEDIAN	PCT MED	
slolpai/slollm-5m		1	4503.57T		4503574627370490	4503.57T	4503.57T	4503.57T	\$0.00	\$0.00	0.00%
slollm-5m		1	3860.21T		3860206823460420	3860.21T	3860.21T	3860.21T	\$0.00	\$0.00	0.00%
slolpai/slollm:5m		1	643.37T		643367803910070	643.37T	643.37T	643.37T	\$0.00	\$0.00	0.00%
claude-opus-4		14	16.21T		16205178862689	3.23M	793.9K	19.09M	\$243.18M	\$10.78	0.18%
gpt-5.5		970	7.20T		7196366570680	1.15B	181.73M	5.46B	\$16.90M	\$1.14K	25.45%
unknown		474	6.50T		6504208521404	2.16M	93.6K	25.33M	\$16.83M	\$0.00	0.00%
gpt-5.4		1091	6.29T		6287426583285	471.02M	80.25M	2.17B	\$4.07M	\$274.59	6.99%
claude-opus-4-7		803	3.18T		3177021574745	840.12M	89.01M	3.60B	\$5.49M	\$747.93	11.56%
claude-opus-4-8		683	3.08T		3077799532554	1.29B	169.57M	5.19B	\$4.10M	\$1.18K	19.29%
claude-opus-4-6		1043	2.12T		2119208240251	441.22M	44.31M	1.89B	\$3.79M	\$394.26	6.79%
gpt-5.3-codex		1006	1.42T		1418748781484	211.94M	29.61M	750.50M	\$497.89K	\$79.37	1.51%
claude-sonnet-4-6		1179	981.32B		981324119143	99.35M	18.23M	441.00M	\$1.77M	\$71.96	1.46%
claude-haiku-4-5		1201	719.81B		719807241043	88.91M	15.78M	339.23M	\$351.42K	\$21.11	0.43%
gpt-5.6-sol		236	572.63B		572630358314	478.31M	80.74M	2.16B	\$402.59K	\$372.23	2.36%
claude-fable-5		494	510.30B		510299584351	357.79M	79.09M	1.23B	\$1.14M	\$689.07	6.15%
gpt-5.2-codex		761	459.04B		459040318672	34.40M	4.14M	182.06M	\$176.84K	\$16.68	0.31%
claude-opus-4-7-20250805		1	449.97B		449972559196	449.97B	449.97B	449.97B	\$4.50M	\$4.50M	100.00%
claude-opus-4-5		685	416.88B		416884580375	59.24M	7.68M	260.88M	\$666.15K	\$67.05	1.11%
gpt-5-codex		404	409.22B		409216640540	61.76M	6.70M	322.21M	\$125.03K	\$20.77	0.36%
codex-mini-latest		7	255.69B		255686835295	50.5K	28.3K	466.1K	\$132.99K	\$0.07	0.01%
deepseek-v4-pro		267	232.12B		232122779249	59.36M	9.97M	231.88M	\$242.57K	\$2.68	0.08%
glm-5.1		293	200.78B		200780489356	43.34M	8.56M	184.14K	\$13.10	\$13.10	0.23%
gpt-5.2		601	195.07B		195072865907	22.87M	3.20M	139.85M	\$74.84K	\$12.02	0.20%
gpt-5.1-codex-max		404	175.90B		175899203415	20.54M	2.94M	83.99M	\$52.99K	\$7.36	0.10%
claude-sonnet-5		300	162.79B		162790673895	98.36M	22.02M	397.27M	\$50.46K	\$39.44	0.35%
gpt-5.4-mini		706	152.42B		152423546159	19.56M	3.36M	108.41M	\$40.75K	\$4.64	0.07%
claude-opus-4-5-20251101		178	121.06B		121059940271	158.59M	27.61M	675.85M	\$104.13K	\$150.34	4.81%
cc-switch-rollup		1	114.52B		114520011575	114.52B	114.52B	114.52B	\$16.98K	\$16.98K	100.00%
claude-sonnet-4-5		772	108.81B		108810671774	20.92M	2.72M	92.54M	\$88.38K	\$15.49	0.25%
gpt-5.3-codex-spark		423	93.63B		93627683705	55.97M	9.89M	202.12M	\$41.53K	\$25.86	0.20%
gemini-3-flash-preview		735	93.07B		9307283166	7.76M	740.0K	51.94M	\$29.77K	\$1.76	0.20%
gpt-5		221	92.27B		92271362650	9.61M	873.2K	44.66M	\$84.92K	\$4.97	0.06%
gpt-5.5-fast		77	82.84B		8284582924	64.73M	2.53M	565.92M	\$90.57K	\$72.10	0.68%
deepseek-v4-flash		219	76.94B		76936573918	35.93M	3.50M	201.75M	\$1.92M	\$0.78	0.02%
glm-5.2		190	75.08B		75083237985	62.51M	13.60M	275.48M	\$44.32K	\$18.80	0.25%
glm-4.7		263	72.62B		72615553550	32.29M	4.20M	156.42M	\$17.41K	\$4.69	0.10%
auto		202	71.78B		71777801371	15.57M	1.15M	158.60M	\$5.92K	\$1.81	0.03%
minimax-m2.7		210	71.09B		71091226390	33.17M	2.61M	182.11M	\$6.11K	\$2.13	0.04%
kimi-for-coding		155	67.86B		67856517157	37.96M	2.63M	209.83M	\$12.18K	\$5.11	0.10%
minimax-m2.7-highspeed		44	63.64B		63638574429	109.05M	10.25M	985.63M	\$8.35K	\$5.22	0.18%
minimax-m3		123	57.01B		57014340122	56.40M	2.97M	278.34M	\$3.99K	\$2.53	0.06%
opal-alpha		1	54.69B		54688576166	54.69B	54.69B	54.69B	\$0.00	\$0.00	0.00%
codex-auto-review		305	53.85B		53854931586	12.05M	1.65M	81.08M	\$2.52K	\$0.78	0.01%
grok-code		201	50.30B		50303926376	4.45M	597.7K	38.06M	\$822.15	\$0.15	0.00%
qwen-coder		1	49.27B		49265147673	49.27B	49.27B	49.27B	\$26.47K	\$26.47K	0.31%
gpt-4o		59	48.92B		48915610816	218.2K	51.2K	1.54M	\$193.37K	\$0.35	0.01%
gemini-2.5-flash		436	47.46B		47455783408	333.1K	51.0K	3.93M	\$30.35K	\$0.07	0.00%
gpt-5.5-xhigh		7	45.05B		45054776487	625.09M	221.22M	5.48B	\$149.84K	\$1.23K	20.08%
mimo-v2.5-pro		132	44.43B		44428262660	44.23M	5.84M	231.05M	\$226.01K	\$1.55	0.05%
gpt-5.1-codex		260	42.78B		42782535978	13.11M	2.25M	77.74M	\$13.12K	\$5.81	0.07%
glm-5		224	42.29B		42290463401	16.18M	1.13M	100.86M	\$25.76K	\$5.60	0.14%
kimi-k2.6		175	38.24B		3823964911	42.14M	7.63M	194.77M	\$739.10K	\$13.96	0.26%
claude-4-5-opus-high-thinking		74	38.23B		38227282515	56.95M	13.24M	228.59M	\$35.59K	\$58.32	1.36%
gpt-5.6-terra		150	33.53B		33530236037	21.94M	6.26M	148.02M	\$13.63K	\$11.25	0.12%
kimi-k2.5		196	31.92B		31919521244	16.03M	1.48M	113.51M	\$10.14K	\$3.28	0.12%
gpt-5.3-codex-xhigh		13	31.69B		31688743086	18.20M	5.01M	366.16M	\$5.35K	\$9.92	0.23%
openai/gpt-oas-20b		14	31.46B		31460376230	202.1K	46.2K	905.3K	\$0.21	\$0.00	0.00%
gemini-3-pro-preview		545	30.51B		30507806121	7.35M	803.1K	35.15M	\$23.73K	\$5.94	0.09%
mimo-v2-5-0		1	29.67B		29669607346	29.67B	29.67B	29.67B	\$296.45	\$296.45	10.07%
gpt-5.5-codex-high		1	26.38B		26377727999	26.38B	26.38B	26.38B	\$23.77K	\$23.77K	7.23%
gpt-5-codex-mini		39	24.39B		24391867278	5.84M	1.60M	12.79M	\$7.19K	\$1.45	0.01%
k2p5		106	24.03B		24028675438	53.77M	8.92M	167.21M	\$4.14K	\$5.04	0.07%
claude-sonnet-4-5-20250929		171	22.20B		22196557374	21.17M	2.71M	83.95M	\$13.35K	\$14.92	0.49%
glm-4.6		80	21.01B		21013746683	6.45M	3.04M	44.53M	\$10.54K	\$1.88	0.04%
gpt-5.1-codex-mini		382	19.46B		19460106466	6.24M	517.1K	37.75M	\$1.49K	\$0.59	0.01%
glm-5-turbo		92	18.63B		18629612413	11.32M	2.40M	90.96M	\$9.81K	\$3.14	0.17%
deepseek-v4-flash-free		148	18.60B		18602873365	25.77M	821.4K	135.13M	\$61.52	\$0.00	0.00%
gpt-5.4-mini-fast		100	18.20B		18196683884	15.51M	2.00M	116.99M	\$4.29K	\$3.53	0.04%
accounts/fireworks/routers/kimi-k2p6-turbo		12	18.18B		18175454835	510.54M	109.36M	799.21M	\$1.60K	\$37.97	0.05%
deepseek-v3		5	17.80B		17797924364	78.4K	40.3K	298.2K	\$7.82K	\$0.01	0.00%
gemini-claude-opus-4-5-thinking		39	17.59B		17594072639	58.44M	13.05M	335.34M	\$80.40K	\$255.94	2.67%
claude-4-6-opus-high-thinking		71	17.01B		17012648594	49.05M	14.31M	173.50M	\$15.97K	\$51.11	1.36%
claude-opus-4-6-thinking		83	16.72B		16721341960	9.83M	1.32M	43.58M	\$43.14K	\$21.83	0.33%
kimi-k2.7-code		71	16.47B		16470613047	24.94M	2.85M	70.35M	\$11.61K	\$4.59	0.12%
minimax-m2.5		151	16.12B		16117981952	11.18M	1.35M	61.42M	\$1.55K	\$0.73	0.01%
claude-opus-4-5-thinking		93	16.10B		16102874355	16.72M	1.86M	111.05M	\$45.94K	\$30.93	0.73%
claude-haiku-4-5-20251001		184	16.10B		16100009936	9.96M	1.46M	58.61M	\$3.77K	\$2.95	0.14%
mimo-v2.5		60	15.33B		15334467126	17.50M	2.20M	103.98M	\$1.85K	\$0.82	0.01%
claude-4-5-sonnet-thinking		69	15.28B		15284383554	57.62M	13.61M	188.42M	\$14.15K	\$50.01	0.76%
big-pickle		534	14.59B		14594450107	954.3K	136.2K	9.89M	\$3.07K	\$0.19	0.00%
gpt-5.4-fast		31	14.02B		14018643204	45.59M	4.42M	571.42M	\$7.72K	\$17.77	0.12%
claude-opus-4-5-high		7	13.96B		13956623578	1.08B	142.88M	2.22B	\$20.25K	\$1.32K	1.64%
gemini-flash-lite-latest		6	13.69B		13685225978	5.27M	421.5K	17.08M	\$357.75	\$0.23	0.01%
gemini-3.1-pro-preview		430	13.23B		13232458789	5.08M	566.4K	24.71M	\$9.28K	\$4.35	0.07%
minimax-m2.5-free		360	13.04B		13037834008	4.58M	775.8K	24.05M	\$393.05	\$0.06	0.00%
composer-2.5-fast		54	12.00B		11998700507	27.35M	4.05M	123.75M	\$5.92K	\$12.06	0.19%
accounts/fireworks/routers/kimi-k2p5-turbo		31	11.93B		11926475628	28.24M	4.18M	375.44M	\$1.47K	\$0.00	0.00%
qwen3.5-plus		69	11.84B		11843405521	7.41M	882.3K	73.30M	\$2.61K	\$0.94	0.02%
kimi-k2-instruct-0905		7	11.63B		11633997341	367.1K	17.7K	6.80M	\$7.11K	\$0.01	0.00%
k2p6		42	11.42B		11420217719	66.85M	27.48M	399.15M	\$1.58K	\$7.51	0.04%
grok-code-fast-1		167	11.41B		11406668335	4.08M	502.2K	28.88M	\$583.19	\$0.33	0.00%
router_client/codex/gpt-5.5		1	11.09B		11089586966	11.09B	11.09B	11.09B	\$0.00	\$0.00	0.00%
claude-4-sonnet		64	10.83B		10834136648	16.13M	2.49M	86.11M	\$5.79K	\$11.23	0.20%
anthropic/claude-opus-4-6		40	10.27B		10268415781	4.83M	1.09M	23.41M	\$14.31K	\$9.42	0.15%
gpt-4-turbo		2	10.16B		10163970046	5.08B	2.54B	7.62B	\$134.22K	\$67.11K	0.80%
composer-2-fast		63	10.16B		10157264158	32.50M	4.33M	103.58M	\$5.16K	\$17.96	0.32%
antigravity-gemini-3-flash		170	10.15B		10146238256						